

Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations

Anonymous Author(s)

Abstract

On a price chart, stablecoin arbitrage during a depeg looks abundant. In practice almost none of it can be traded at a profit. During the March 2023 USDC depeg, 34.3% of one-minute windows showed a cross-venue price basis above 10 bps, yet only one in twelve of those apparent opportunities survived VWAP slippage, order-book depth depletion, and settlement latency (block bootstrap 95 % CI: $8\times-24\times$). We introduce Stablecoin StressBench to quantify three compounding gaps between an apparent arbitrage opportunity and a captured one. The first gap is execution: price labels overstate opportunities by $12\times$. The second gap is prediction: even with execution-aware labels, every calm-trained model loses money and the hindsight oracle stays more than 260 bps ahead of the best deployable model. The third gap is policy: cross-mechanism meta-labeling trained on Terra/LUNA achieves +82.5 bps on SVB (50.8% of oracle), while a conditioned RL agent at identical training density earns -29.2 bps, isolating explicit binary supervision as the requirement that reward optimisation cannot replicate at low positive rates.

Keywords

stablecoin arbitrage, market microstructure, execution-aware labels, financial benchmarks, cryptocurrency

ACM Reference Format:

Anonymous Author(s). 2026. Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations. In *Proceedings of ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 8 pages.

1 Introduction

When Silicon Valley Bank collapsed in March 2023, USDC briefly depegged and cross-venue price gaps appeared profitable on every chart. Almost none of them could be traded at a profit. A price chart shows twelve times more stablecoin arbitrage than a trader can capture, and this paper traces where the other eleven twelfths disappear. VWAP slippage consumed the spread as orders cleared depth, taker fees eroded the margin, and CEX-to-CEX settlement latency turned a profitable window into a loss before the transfer confirmed. Only 2.88% of one-minute windows survived all three costs and remained net-profitable. We call dislocations that appear on price charts but vanish under execution *optical*, and the surviving fraction *executable*. Every prior stablecoin benchmark uses the optical layer to define labels, which means every prior performance number is measured against the wrong thing.

Stablecoin StressBench corrects this by constructing per-minute execution-aware labels for the USDC/SVB event using a full VWAP book-walking model applied to Binance L2 depth, separating three layers: *optical* (price basis above threshold), *executable* (net profit

after VWAP, fees, and settlement), and *predictable* (correctly identified ex ante). The distance between each pair of layers is a distinct, measurable gap, and the benchmark is designed to quantify all three.

Stablecoin StressBench measures three compounding gaps that separate an apparent arbitrage opportunity from a captured one. The first gap is execution. A price chart shows a dislocation that an order book cannot fill at a profit, and we measure this gap at twelve to one. The second gap is prediction. Even when the labels are correct, every model trained on calm data loses money, and the hindsight oracle stays more than 260 bps ahead of the best deployable model. The third gap is policy. A learning agent can be given the right training density and still fail to trade profitably, because reward optimisation does not supply the binary supervision that meta-labeling provides. The remainder of the paper measures each gap in turn, in §6.1, §6.2, and §6.3.

Contributions.

- (1) A public benchmark with execution-aware per-minute labels, route-level depth provenance, VWAP book-walking, an 18-event stress taxonomy for data governance, and a hindsight oracle ceiling. (§3–5)
- (2) The first evidence, on the single cross-mechanism pair with route-complete test data, of order-book execution microstructure transfer across structurally distinct stress mechanisms: meta-labeling trained on Terra/LUNA achieves +82.5 bps (50.8% oracle capture) on the SVB test split, while every calm-trained model loses money. (§6, Table 12)
- (3) An RL diagnostic showing that positive-label density is necessary but not sufficient: a conditioned PPO-GRU at 17% positive density earns -29.2 bps versus meta-labeling's +82.5 bps at identical density. (§6, Table 15)

2 Related Work

Stablecoin StressBench sits at the intersection of three bodies of work.

Research on stablecoin stress events motivates the benchmark. Uhlig [1] analyses the Terra/LUNA algorithmic collapse, which we use as an out-of-distribution training split precisely because it is mechanistically different from the SVB bank shock. Griffin and Shams [2] show that quoted prices can diverge from fundamentals, motivating our distinction between optical and executable dislocations. Hernandez Cruz et al. [21] study DEX liquidity during SVB, but stop short of execution-layer modelling. None of this prior work provides per-minute executable labels or a hindsight oracle ceiling.

The limits-to-arbitrage literature provides the economic framing that our labels operationalise. Shleifer and Vishny [4] and Gromb and Vayanos [5] establish that execution frictions bound arbitrage even when price gaps exist. Makarov and Schoar [7] document that cross-venue crypto gaps vanish once depth and latency are modelled, the same mechanism behind our $12\times$ ratio. Hautsch et

Table 1: Route-leg depth provenance.

Leg	SVB Mar 2023	2024 windows	
Buy BTC-USDC	kline-proxy [†]	real L2 (futures)	†BTC-USDC
Cross USDC-USDT	real L2 (Mar 12–20)	real L2 (futures)	
Sell BTC-USDT	real L2 (futures)	real L2 (futures)	
Ref BTC-USD	candle-proxy	candle-proxy	

perpetual listed Jan 2024; all 2022 windows kline-proxy.

al. [3] quantify settlement latency as a first-order cost, which we implement as the $\delta=5$ bps penalty; Cont and Kukanov [6] establish VWAP price as the correct profitability object. Gogol et al. [18] measure a Maximal Arbitrage Value for AMM-to-CEX routes; our oracle extends this idea to stress events, route-level CEX simulation, and model evaluation. Where prior theoretical work characterises frictions, we measure them per-minute with empirical depth data.

Financial ML and benchmarks form the methodological context. López de Prado [8] introduces meta-labeling as a framework for filtering primary signals with a learned secondary classifier. We extend it to the cross-mechanism setting, training on Terra/LUNA and evaluating on SVB, which to our knowledge has not been done. Cintra and Holloway [9] demonstrate cross-event transfer using AMM reserve depletion signals, but their signals are not CEX-applicable, no oracle evaluation is provided, and no execution-aware labels are constructed. On RL execution, Moallemi and Wang [16] and Ning et al. [17] establish the finite-horizon MDP framing we adopt; Mohl et al. [22] show GPU-scale IPPO on 400 million real orders; and Olby et al. [23] use a reactive Simudyne simulator to benchmark timing policies against an Almgren–Chriss frontier. These works show RL succeeds in reactive environments with frequent positive labels. Our diagnostic establishes the preconditions that must hold before reactive training can succeed.

3 Benchmark Design

3.1 Data Sources and Instruments

The final dataset contains 56,134 1-minute bars across six windows (Table 2). Price features are sourced from Binance Vision (klines and aggTrades) and Coinbase REST candles. Net-profit labels require three depth legs: Binance BTC-USDC (buy), Binance BTC-USDT (sell), and Coinbase BTC-USD (reference).

Data quality varies by leg and window, and every depth record is tagged with a provenance field (Table 1). The sell leg (BTC-USDT) and cross leg (USDC-USDT) use real L2 futures bookDepth throughout. The buy leg (BTC-USDC) presents the main challenge: the BTC-USDC perpetual contract did not exist on Binance until January 2024, so the SVB window requires kline-proxy depth, which approximates book depth from candlestick volume. For all 2024 windows, real L2 data is available. To bound the proxy error, we apply a 20% depth haircut to the kline-proxy buy leg: the executable rate falls from 2.88% to 2.40%, the ratio widens from 12× to 14×, and the oracle remains above +130 bps. The headline result is therefore robust to the approximation.

3.2 Event-Based Split Design

Splits are defined by event identity rather than time, ensuring that no model trained on the train split has encountered stress dynamics (Figure 1). Threshold calibration uses the Terra/LUNA validation split, which is an independent stress event, so the SVB test split is

Table 2: Dataset splits (56,134 1-min bars).

Split	Event	Min.	Pos.
Train	Calm control (×3)	28,776	0.00%
Validation	Terra/LUNA May 2022	11,526	2.30%
Test	USDC/SVB Mar 2023	15,832	2.88%




Figure 1: Cross-quote basis (bps) during the SVB test window (Mar 10–20 2023). USDC basis (navy): 12.4% of minutes > 10 bps; USDT basis (red): 27.4%; max-abs (gold): 34.3%. The price-to-execution gap is 12×.

never observed during model development. Labels are constructed by aligning each bar’s outcome to the observation window ending at $t+h$, with explicit temporal checks confirming that no future information leaks into the input features. The 2.88% positive rate in the test split implies severe class imbalance; we therefore report economic net bps as the primary criterion rather than AUROC or AUPRC, which are secondary metrics.

3.3 Historical Stress Taxonomy: Data-Tier Governance

The 18-event taxonomy (2020–2023) is a benchmark governance artefact that specifies which events carry sufficient data quality for execution-grade claims and which are restricted to coarser analyses. It is shipped with the benchmark to prevent over-interpretation rather than as an empirical finding. Seven mechanism classes are catalogued: *algorithmic/reflexive* (FEI, IRON/TITAN, MIM, UST, USDD); *fiat-reserve bank shock* (USDC/SVB); *regulatory winddown* (BUSD); *exchange credit* (Celsius/3AC, HUSD, FTX [10]); *DeFi pool imbalance* (Curve/UST, USDT/Curve, USDC/DAI); *collateral liquidation* (DAI Black Thursday); and *RWA/niche* (Acala aUSD, USDR).

Data availability governs which claims are permitted (Table 3). The Tier-A event (USDC/SVB test split) supports execution-gap, oracle-gap, and model evaluation claims because route-complete L2 depth is available for all three arbitrage legs. Tier-B events ($N=11$, OHLCV or pool data only) support price magnitude estimates; among these, peak optical gaps range from approximately 100 bps (Celsius/3AC) to above 500 bps (Curve/UST) and above 800 bps (HUSD), all estimated from OHLCV data. Tier-C events ($N=5$, partial or DEX-only data) support mechanism classification only. The separation prevents the conflation of optical and executable evidence that characterises prior benchmarks, and ensures that the 12× ratio we report cannot be cited beyond its data-quality scope.

Table 3: Data tiers, mechanisms, and execution frictions.

Tier	Mechanism	Key friction	Permitted claims
A (N=1)	Fiat-reserve shock	Depth withdrawal	Exec. gap; oracle gap; P&L
B (N=11)	Alg., exchange, DeFi	AMM venue seg.	slippage; Price magnitude (est.)
C (N=5)	RWA, niche	Thin market	Taxonomy only

Table 4: Comparison with related benchmarks and datasets.

✓ = fully provided; ~ = partially; – = absent.

Work	Stress event	Exec. labels	L2 prov.	Oracle	Cross-event
Cintra & Holloway [9]	✓	–	AMM	–	✓
Hernandez Cruz et al. [21]	✓	–	DEX	–	–
Gogol et al. [18]	–	~	CEX	MAV	–
Adams et al. [19]	–	~	CEX	–	–
StressBench (ours)	✓	✓	CEX	✓	✓

3.4 Benchmark Positioning

Table 4 positions StressBench against the closest prior work along five axes. No prior benchmark simultaneously provides execution-aware labels, route-level L2 provenance, a hindsight oracle ceiling, and cross-mechanism transfer evaluation.

3.5 Benchmark Versioning

StressBench v1.0 ships with the SVB test split. Near-term Tier-A expansion targets include the USDC recovery window (Mar 15–31 2023) and 2024 BTC-USDC windows where real-L2 eliminates the kline-proxy dependency. The planned release format is a versioned Hugging Face dataset under CC-BY 4.0.

3.6 Feature Sets

Four nested feature sets isolate the marginal information value of

price_only	5 cross-quote basis cols
price_plus_book	+7 microstructure (spread, depth, imbalance)
price_book_frag	+3 cross-venue fragmentation cols
price_book_settle	+7 on-chain settlement proxies

4 Execution-Aware Labels

Labels operationalise the three-layer framework from §1. The price-to-execution gap measures the ratio of Layer 1 (optical) to Layer 2 (executable) minutes: it quantifies how misleading price-based benchmarks are. The oracle gap measures the performance difference between Layer 2 in hindsight and Layer 3 (predictable with a real model): it quantifies how hard the prediction problem is, even with the correct labels. Both gaps are nonzero, independent, and large. This section details how they are constructed.

4.1 Cross-Quote Basis and Labels

$$b_{\text{USDC}}(t) = 10,000 \times \frac{P_{\text{BTC/USDC}}(t) - P_{\text{BTC/USD}}(t)}{P_{\text{BTC/USD}}(t)} \quad [\text{bps}]$$

$b_{\text{max}} = \max(|b_{\text{USDC}}|, |b_{\text{USDT}}|)$. Basis label: $y_{\text{basis}}(t) = 1[|b_{\text{USDC}}(t+h)| > \tau]$. Primary task: basis_usdc_1m_gt10bps.

4.2 VWAP Execution and Net Profit Labels

$$\text{net}(q, t) = 10,000 \times \left[\frac{\text{VWAP}_{\text{sell}} - \text{VWAP}_{\text{buy}}}{P_{\text{ref}}} - f_{\text{buy}} - f_{\text{sell}} - \delta_{\text{settle}} \right] \quad [\text{bps}] \quad (1)$$

$$y_{\text{arb}}(q, h, t) = 1 \left[\max_{s \in [t, t+h]} \text{net}(q, s) > c \right] \quad (2)$$

Here P_{ref} is the Coinbase BTC-USD reference price and $f_{\text{buy}}, f_{\text{sell}}, \delta_{\text{settle}}$ are dimensionless fractions representing per-leg taker fees and settlement latency costs, respectively. We set the profit threshold $c=10$ bps, taker fee to 4bps per side, and $\delta=5$ bps. The primary economic task is executable_arb_q10000_5m. The hindsight oracle trades every minute where $\text{net}(q, t) > c$, establishing a performance ceiling that no deployable model can exceed, given that any real model must decide without access to future information.

5 Models and Evaluation

5.1 Model Ladder

The ladder is constructed so that each rung tests a distinct hypothesis about the source of the oracle gap, with failure at each rung ruling out a candidate explanation. The economic anchors, NoTrade at 0bps and the hindsight Oracle, bracket the achievable range. Price-threshold rules ($\text{PriceBasis}\{10, 25\}$ bps, GrossArbThreshold) test whether simple signal filters suffice. Tabular classifiers (logistic regression, Lasso, Random Forest, XGBoost, LightGBM) are trained on four nested feature sets that progressively add microstructure, cross-venue fragmentation, and on-chain settlement signals; their failure tests whether richer features close the gap. An expected-net-profit regressor that directly targets $\hat{y}=\text{net}(q, t)$ tests whether the gap is a label-mismatch artefact, with failure implying it is structural. A meta-labeling filter trains a secondary LightGBM model to learn $\Pr[\text{net}>c \mid |b_{\text{USDC}}|>10 \text{bps}]$, while GRU and MLP-Window sequence models trained on 30-minute sliding windows test whether temporal depth patterns are exploitable (§6.2). The cross-mechanism meta-labeling variant is the same architecture but trained on Terra/LUNA rather than calm data, directly testing whether the training distribution is the binding constraint.

5.2 RL Execution Policy Baseline

We frame the entry-timing problem as a finite-horizon MDP. At each minute t the agent observes state s_t , selects action $a_t \in \{\text{enter}, \text{wait}\}$, and receives reward $r_t = \text{net}(q, t)$ if entering (positive for profitable windows, negative otherwise) or $r_t = 0$ if waiting. The state s_t is the 30-minute look-back window of price_plus_book features, the smallest feature set that delivers meaningful microstructure lift over price-only signals (see §6.3 for SHAP evidence). This feature set is available in real time without settlement-latency dependency. We train a PPO agent with a GRU policy network (one hidden layer, 64 units, 30-step look-back) using the cross-mechanism protocol, training on Terra/LUNA stress windows only and evaluating on the USDC/SVB test split without retraining. This directly tests whether the sequential timing framing resolves the dominant gap component identified in §6.2.

5.3 Calibration and Evaluation

Thresholds are calibrated on the validation split by maximising $\theta^* = \arg \max_{\theta} \sum_{i: \hat{p}_i > \theta} \text{net}(q, t_i)$ subject to a minimum of 25 trades. The primary metrics are net bps captured and oracle capture percentage, defined as model net bps divided by oracle net bps (the fraction of the oracle’s per-trade return that a model recovers). Secondary metrics are AUROC and AUPRC. NoTrade at 0 bps is the economic floor.

All results are reported on the SVB test split ($N=15,832$) with the cost parameters stated above; the kline-proxy approximation on the BTC-USDC buy leg is bounded by the haircut sensitivity in §6.1. Five explicit checks guard against leakage and evaluation artefacts. Random label shuffles produce zero or negative P&L for all models. Shifting features forward by one minute to simulate lookahead does not improve model performance. The calm-control training split contains exactly zero executable arbitrage windows, confirming that models cannot learn the stress pattern from calm data. Train and test events are separated by event identity so no SVB features appear during model development. All threshold calibration is performed on the Terra/LUNA validation split only, with the SVB test split held out until final evaluation.

Benchmark Card. Table 5 summarises the five evaluation tasks shipped with StressBench v1.0.

Table 5: StressBench v1.0 benchmark card.

Task	Train	Val.	Test	Primary metric
optical_basis	calm	Terra	SVB	AUROC, AUPRC
exec_arb_q10k	calm	Terra	SVB	net bps, oracle cap.
exec_arb_q50k	calm	Terra	SVB	net bps
cross_mech_transfer	Terra	Terra	SVB	net bps, oracle cap.
policy_entry (RL)	Terra	Terra	SVB	net bps, trades

6 Results

6.1 Gap 1: Execution

The first gap is between what a price chart shows and what an order book will fill. At a 10 bps threshold, 34.3% of test minutes show a dislocation and only 2.88% remain profitable after VWAP slippage, taker fees, and settlement latency. The ratio is twelve to one, and it widens as trade size grows (Table 6). This ratio is not a threshold artefact: it persists across all tested basis levels and is robust to data-quality concerns. Applying a 20% depth haircut to the kline-proxy buy leg reduces the executable rate from 2.88% to 2.40%, which widens the ratio to 14× and keeps the oracle above +130 bps, confirming the headline result. The gap also deepens with institutional scale. Figure 3 shows the executable rate falling from 2.88% at \$10K to 0.03% at \$500K, with the steepest drop between \$10K and \$50K. Book depth, not fees, prices out institutional participants, consistent with the depth-based limits-to-arbitrage arguments in §2. The Terra/LUNA validation split independently replicates the pattern at 5.9× (Table 14), providing directional evidence that the gap is not SVB-specific.

Table 7 shows that the gap is robust across the full cost-parameter space: the executable rate varies only between 2.84% and 3.66% as

Table 6: Price-to-execution gap (\$10K, 1-min horizon, test split).

Thresh.	$P_{\max}$	P_{USDC}	Exec.	Ratio
0 bps	95.86%	82.96%	3.34%	29×
5 bps	61.99%	25.32%	3.03%	20×
10 bps	34.33%	12.45%	2.88%	12×
25 bps	9.54%	6.48%	2.62%	3.6×




Figure 2: Oracle gap decomposition (\$10K notional). Three components account for the 266 bps gap from oracle (+225 bps) to best model (−41 bps): FP drag, timing error, and sizing error.




Figure 3: Notional scaling (SVB test split). The executable rate collapses from 2.88% at \$10K to 0.03% at \$500K: book depth, not fees, is the binding institutional constraint.

fee and settlement assumptions are swept, and the ratio to the optical rate remains above 3.5× throughout. Claim boundaries and data-tier scope are stated in §5.3.

6.2 Gap 2: Prediction

The second gap is between a profitable window and a model that can find it in advance. Even with execution-aware labels, every model trained on calm data loses money on the executable task. The hindsight oracle earns +225.0 bps while the best deployable model reaches −41.1 bps, a structural gap of more than 260 bps that richer features and sequence models do not close (Table 8; Figure 4).

Table 7: Executable rate (%) under fee and settlement variation (10 bps threshold, 1-min horizon, \$10K notional, SVB test split). Taker fees: low = 2 bps, base = 4 bps, high = 8 bps, inst. = 10 bps per side.

Taker fee	Settlement penalty (bps)			
	0	2	5	10
Low (2 bps)	3.50	3.34	3.16	2.93
Base (4 bps)	3.34	3.20	3.03	2.88
High (8 bps)	3.20	3.08	2.96	2.84
Inst. (10 bps)	3.66	3.42	3.20	2.96

Table 8: Oracle gap summary (USDC/SVB test split). *Only non-oracle tabular result above zero. ‡GRU sequence model, 30-min window.

Task / Model	Oracle	Best model	Gap	AUPRC
basis_usdc_1m	+162.2 bps	+17.2 bps*	145 bps	0.253
exec_arb_q10k_5m	+225.0 bps	−41.1 bps	266 bps	0.060
exec_arb_q50k_5m	+147.1 bps	−113.0 bps	260 bps	0.063
GRU (seq)‡, basis_usdc_1m	+162.2 bps	−239.0 bps	401 bps	—

Kline-proxy buy leg; 20% haircut sensitivity in \$6.1.




Figure 4: Running mean net bps per trade (USDC/SVB test split). Oracle (gold) +162.2 bps; cross-mech meta-label (green) +82.5 bps; LightGBM (navy) +17.2 bps; PriceBasis10bps (red) −269.5 bps; NoTrade (grey dashed) 0 bps. Shaded region: SVB shock window.

Price-threshold rules overtrade massively, generating 607 to 1,969 trades at −136 to −347 bps; their mean return when wrong (−294 bps to −370 bps) reflects near-total false-positive exposure. LightGBM (+17.2 bps, 106 trades, AUROC 0.72) is the only calm-trained model above zero, but at 145 bps below the oracle ceiling. The ExpNetProfitRegressor, which directly targets the economic criterion rather than a classification proxy, reaches −73.8 bps, ruling out label mismatch as the explanation: the gap is structural (Table 9).

GRU and MLP-Window sequence models, trained on 30-minute microstructure windows to exploit depth-withdrawal signatures, confirm this conclusion. The GRU achieves AUROC 0.80 yet earns −239 bps, and the MLP-Window −341 bps (Table 10). High ranking accuracy does not translate to positive returns when 97% of *all* test windows are non-executable and the training split contains zero positive stress examples.




Figure 5: Precision-recall curves (basis_usdc_1m_gt10bps, SVB test split). PriceBasis achieves AP 0.55 and AUROC 0.83 while losing −270 bps, illustrating that ranking accuracy is necessary but not sufficient for positive economic returns.

Table 9: Tabular model ladder (basis_usdc_1m_gt10bps). Hit%: fraction of trades net-profitable. † Ceiling. *Only result above zero.

Model	Feat.	Netbps	Trades	AUROC	Hit%
Oracle†	—	+162.2	315	—	100%
lgbm*	all	+17.2	106	0.717	33%
logistic	all	−75.8	648	0.143	0%
gross_arb	px	−136.4	607	0.531	4%
price_10	px	−270.0	1969	0.833	6%
price_25	px	−347.3	1025	0.746	5%
no_trade	—	0.0	0	—	—
ExpNetProfit	px	−73.8	537	—	3%

all=price_book_settle.

Table 10: Sequence model results (basis_usdc_1m_gt10bps).

Model	Netbps	Trades	AUROC	Gap
GRU (30-min)	−239.0	656	0.798	401 bps
MLP-Window (30-min)	−340.7	1,052	0.648	502 bps

Calm-control trained; threshold calibrated on validation.

Table 11 converts the model ladder into an economic confusion matrix, separating the cost of false positives from the cost of missed oracle trades. Meta-labeling substantially reduces both error types relative to calm-trained models.

Figure 2 decomposes the 260 bps gap into three measurable components. False-positive cost is the dominant and directly quantifiable drag: PriceBasis10bps produces 1,969 trades of which 1,851 are false positives. Of those, 581 are attributable specifically to route-direction mismatch (USDT dislocations triggering on the max-abs basis when the USDC route is not simultaneously profitable) at −38.7 bps each; the remainder reflects imprecise timing. The remaining gap splits into timing error (estimated 40–80 bps), arising from entering 1–2 minutes late in a declining depth regime, and sizing error (40–100 bps), since depth depletion reduces the profitable fraction from 2.88% at \$10K to 1.62% at \$50K.

Table 11: Economic confusion matrix (basis task, SVB test split). Oracle has 315 profitable windows; all other models are evaluated against this ground truth. FP loss and FN missed profit are cumulative bps across the test split.

Model	TP	FP	FN	Net bps
Oracle	315	0	0	+162.2
Meta-label (Terra) [†]	~202	~195	~113	+82.5
LightGBM (calm)	35	71	280	+17.2
PriceBasis10bps (calm)	118	1851	197	-270.0
NoTrade	0	0	315	0.0

FN = 315 – TP. [†]Meta-label TP/FP are approximate: Hit% is not in Table 9; the ~202

TP estimate uses oracle-capture rate (50.8%) as a proxy for hit rate, which overstates

TP if profitable trades earn less than oracle bps per trade.

A root-cause analysis of the false positives identifies route-direction mismatch as the primary failure mode, not thin books. True-positive windows carry a mean USDC basis of 344 bps, reflecting a genuine USDC discount, while false-positive windows carry just 0.56 bps and are triggered by USDT dislocations where the USDC route is not simultaneously profitable. Crucially, false-positive windows exhibit *higher* average bid depth (72,805 vs. 59,961 BTC units) than true positives. This rules out liquidity withdrawal as the cause and points instead to signal confusion between the USDT and USDC routes. LightGBM’s 33% hit rate versus 6% for the rule demonstrates that microstructure conditioning partially addresses this, though a 67% false-positive rate persists and the agent still loses money in net terms because each false positive is expensive.

6.3 Gap 3: Policy

The third gap is between a correct signal and a profitable policy. Cross-mechanism meta-labeling trained on Terra/LUNA is the only method that crosses zero, earning +82.5 bps and recovering half of the oracle return. A reinforcement learning agent given the same 17% training density earns -29.2 bps, which isolates explicit binary supervision as the ingredient that reward optimisation cannot replicate at low positive rates.

6.3.1 Cross-Mechanism Transfer. Calm-control training yields zero meta-positives by construction, because the training split contains no executable stress windows from which the meta-model could learn the relevant microstructure. Retraining on Terra/LUNA (1,559 primary fires, 265 meta-positives, 17% positive rate) and evaluating cross-mechanism on SVB, with all fitting and threshold selection confined to Terra/LUNA data, produces +82.5 bps (397 trades, 50.8% oracle capture) versus zero with calm training (Table 12).

Table 12: Cross-mechanism meta-labeling results (basis_usdc_1m_gt10bps, price_plus_book).

Training split	Net bps	Trades	Oracle cap.
Oracle (ceiling)	+162.2	315	100.0%
Terra/LUNA (cross-mech.)	+82.5	397	50.8%
Calm-control (baseline)	0.0	0	0.0%

signal $|b_{\text{USDC}}| > 10$ bps; threshold calibrated on validation.

The reason transfer works is visible in the SHAP decomposition (Table 13). Price features carry 94.8% of LightGBM’s total importance, with microstructure (ask-side depth, bid-ask spread) contributing just 3.4% combined. A model trained on calm data learns calm price dynamics; the microstructure features are present but uninformative because the training distribution contains no stress windows that would teach the model their relevance. Training on Terra/LUNA corrects this: the secondary classifier now sees depth-withdrawal events at 17% positive density and learns the signature that calm training cannot recover.

Table 13: Top-5 SHAP feature importances (LightGBM, basis_usdc_1m_gt10bps, price_book_settle, SVB test split). Price features dominate at 94.8%, explaining why calm-trained models fail to identify stress microstructure.

Feature	Category	Importance (%)
basis_usdc_bps	Price	49.6
basis_primary_bps	Price	32.8
deviation_1usd_bps	Price	10.1
basis_usdt_bps	Price	2.0
depth_ask_10bp	Book	0.5
All price features	Price	94.8
All book/frag features	Book	3.4
Settlement features	Settle	1.8

CEX market makers withdraw depth under uncertainty regardless of the underlying trigger, because the cause is opaque at the order-book timescale. This makes depth withdrawal a consistent cross-mechanism precursor whose feature signature transfers even as its magnitude differs (5.9× for Terra/LUNA versus 12× for SVB, shown in Table 14).

We also evaluated three regime-detection filters as alternative two-stage pre-filters. CUSUM achieves near-perfect recall (0.9995) but at an 83.5% false alarm rate, making it impractical as a standalone filter. EWMA reaches 86.9% accuracy with 0.95% recall at a 0.65% false alarm rate. BOCPD [15] achieves AUROC 0.229 on CEX cross-quote data, which contrasts with Cintra and Holloway [9], who detected the Terra depeg 12 hours early from Curve AMM pool reserves. The difference is explained by the signal type: AMM reserves accumulate gradually as liquidity providers withdraw, making them suitable for changepoint detection. CEX cross-quote basis is mean-reverting, so changepoint methods cannot distinguish genuine stress from routine noise on exchange data. Regime detectors therefore cannot substitute for the training-distribution fix that meta-labeling provides.

6.3.2 Density Is Necessary and Supervision Is Decisive. Meta-labeling and a conditioned PPO-GRU share identical training density (17% positive rate on primary-signal windows) yet differ by 111.7 bps in net return. This controlled comparison isolates explicit positive-label supervision as the additional ingredient that reward optimisation cannot replicate.

The unconditioned PPO-GRU degenerates to zero entries, as the 2.3% positive rate in the full Terra/LUNA stream is too sparse to generate a reward signal that calibrates entry timing. Conditioning on primary-signal windows resolves the density problem: the agent

Table 14: Cross-mechanism comparison (kline-proxy labels for Terra/LUNA; directional evidence only).

Metric	USDC/SVB (Tier-A)	Terra/LUNA (val.)	
Mechanism	Fiat-reserve	Alg./reflexive	
Price > 10 bps	34.3%	13.5%	[†] Kline-proxy;
Executable	2.88%	2.30%	
P-to-E ratio	12×	5.9×	
Oracle net bps	+225.0	+88 [†]	
directional only.			

Table 15: RL policy versus supervised baselines (basis_usdc_1m_gt10bps, SVB test split).

Model (training)	Net bps	Trades	
Oracle (hindsight ceiling)	+162.2	315	
Meta-labeling (Terra/LUNA*)	+82.5	397	*Cross-mechanism:
PPO-GRU conditioned [†] (Terra/LUNA*)	−29.2	919	
PPO-GRU unconditioned (Terra/LUNA*)	degenerate	0	
GRU supervised (calm-ctrl)	−239.0	656	
MLP-Window supervised (calm-ctrl)	−340.7	1,052	

Terra/LUNA train, SVB eval. [†]Conditioned on $|b_{USDC}| > 10$ bps windows, 17% positive rate.

trades (919 entries, AUROC 0.87) but earns −29.2 bps. The distinction between the conditioned agent and meta-labeling is qualitative. The meta-labeling classifier receives direct binary feedback, where $y=1$ marks exactly the profitable windows, while the RL agent must infer the same subset through a reward signal that is negative on 83% of entries. The supervision is categorically cleaner, not merely denser. The 919-trade overtrading pattern is itself informative. Policies that fail due to insufficient training converge toward NoTrade, not sustained false-positive entry. The 919-trade result is therefore more consistent with a timing-transfer mismatch from Terra/LUNA to SVB, meaning the agent learned Terra/LUNA’s specific depth-withdrawal sequence and fires on it in the SVB context even when the timing does not match. This interpretation has a constructive implication. If the failure reflects a mechanism mismatch rather than an RL capability limit, then multi-mechanism training that exposes the agent to depth withdrawal events across diverse stress triggers should improve selectivity without requiring a fundamentally different RL algorithm. Combining the cross-mechanism protocol established here for supervised transfer with a reactive simulator that provides market-impact feedback is the natural next step.

7 Limitations

The BTC-USDC buy leg uses kline-proxy depth for the SVB window, since the perpetual contract did not exist on Binance until 2024. A 20% depth haircut confirms robustness (Table 7).

We validate the $\delta=5$ bps settlement penalty assumption using on-chain Etherscan data for 100,000 USDC transfers across the SVB window (Table 16). The median ERC-20 transfer cost on 10K notional was 1.4–3.2 bps on nine of ten days, confirming that $\delta=5$ bps is conservative. The exception is March 11 (peak depeg), where gas spiked to 61.5 gwei and the P95 cost reached 8.6 bps, the only day the assumption is tight. Burn events (institutional USDC redemptions to Circle) peaked on March 13 (29 events) and transfer volume hit \$2.4B on that day, providing on-chain corroboration of

the depth depletion visible in the CEX order book during the same window.

Table 16: On-chain settlement cost validation (USDC ERC-20 transfers, 100K sample, Mar 10–20 2023, ETH $\approx$ 1,580). $\delta=5$ bps is conservative on all days except the Mar 11 gas spike at peak depeg.

Date	Med gas (gwei)	Med cost (bps)	P95 cost (bps)	Burns	
Mar 10	30.8	3.16	5.70	4	[†] Only day where
Mar 11	61.5	6.31	8.64 [†]	16	
Mar 12	21.6	2.21	3.68	22	
Mar 13	26.1	2.68	4.47	29	
Mar 14–19	13.9–24	1.43–2.46	2.44–4.11	0–2	

P95 cost exceeds $\delta=5$ bps. Gas data from 100K USDC transfers via Etherscan API.

Tier-A coverage is currently limited to the SVB event; extending to AMM or exchange-credit events requires route-complete L2 archives that are not yet available.

The RL results carry two important caveats. First, the conditioned PPO-GRU’s −29.2 bps result may reflect overfitting to Terra/LUNA’s specific depth-withdrawal timing rather than a fundamental RL limitation, which additional cross-mechanism training pairs could resolve. Second, unlike reactive-simulator approaches such as Olby et al. [23], the PPO-GRU is trained on historical replay without market-impact feedback. In a thin-book stress regime this systematically miscalibrates entry sizing because the agent observes book depth as it was rather than as it would be after its own orders clear liquidity, which is a plausible contributor to the 919-trade overtrading pattern. Integrating the cross-mechanism protocol shown here into a reactive simulator is the most direct path to resolving the timing component of the oracle gap.

8 Conclusion

The 12× optical-to-executable gap establishes that stablecoin arbitrage benchmarks built on price labels measure the wrong quantity. The executable fraction of apparent opportunities is small, shrinks with notional size, and is invisible to every model trained on calm-market data. Cross-mechanism training on a structurally different stress event is the only empirically validated path to positive net returns, and the RL diagnostic provides a precise account of why: label density enables trading, but binary supervision over the profitable subset provides the directional accuracy that reward optimisation cannot replicate at low positive rates. The benchmark, labels, oracle implementation, and training scripts are released under CC-BY 4.0 (URL redacted for review). StressBench is designed to support four lines of future work: execution-aware dislocation detection under cost uncertainty; cross-mechanism stress transfer to new event types as Tier-A coverage expands; RL execution policy learning with the cross-mechanism initialisation protocol established here; and reactive simulation of thin-book stablecoin markets, where the oracle ceiling provides a principled performance target. The benchmark harness additionally ships four practitioner desk-rule baselines (route-consistent, depth-adjusted, spread-depth, and latency-haircut rules) that exercise the same evaluation pipeline and serve as interpretable reference points for future model comparisons.

Acknowledgements

Anonymised for review.

References

- [1] Uhlig, H. (2022). A Luna-tic stablecoin crash. *NBER Working Paper 30256*.
- [2] Griffin, J. M. and Shams, A. (2020). Is Bitcoin really untethered? *Journal of Finance*, 75(4):1913–1964.
- [3] Hautsch, N., Scheuch, C., and Voigt, S. (2018). Building trust takes time: Limits to arbitrage for blockchain-based assets. *arXiv:1812.00595*.
- [4] Shleifer, A. and Vishny, R. W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1):35–55.
- [5] Gromb, D. and Vayanos, D. (2010). Limits of arbitrage: The state of the theory. *Annual Review of Financial Economics*, 2:251–275.
- [6] Cont, R. and Kukanov, A. (2013). Optimal order placement in limit order markets. *Quantitative Finance*, 17(1):21–39.
- [7] Makarov, I. and Schoar, A. (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics*, 135(2):293–319.
- [8] López de Prado, M. (2018). *Advances in Financial Machine Learning*. Wiley.
- [9] Cintra, R. and Holloway, C. (2023). Detecting depegs: Towards safer passive liquidity provision on Curve Finance. *arXiv:2306.10612*.
- [10] Vidal-Tomás, D., Briola, A., and Aste, T. (2023). FTX’s downfall and Binance’s consolidation. *arXiv:2302.11371*.
- [11] Gatev, E., Goetzmann, W. N., and Rouwenhorst, K. G. (2006). Pairs trading: Performance of a relative-value arbitrage rule. *Review of Financial Studies*, 19(3):797–827.
- [12] Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *JRSS-B*, 58(1):267–288.
- [13] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1):5–32.
- [14] Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *KDD*, pp. 785–794.
- [15] Adams, R. P. and MacKay, D. J. C. (2007). Bayesian online changepoint detection. *arXiv:0710.3742*.
- [16] Moallemi, C. C. and Wang, M. (2022). A reinforcement learning approach to optimal execution. *Quantitative Finance*, 22(6):1051–1069.
- [17] Ning, B., Lin, F. H. T., and Jaimungal, S. (2021). Double deep Q-learning for optimal execution. *Applied Mathematical Finance*, 28(4):361–380.
- [18] Gogol, K., Messias, J., Miori, D., Tessone, C., and Livshits, B. (2024). Quantifying arbitrage in automated market makers: An empirical study of Ethereum ZK rollups. In *MARBLE 2024*.
- [19] Adams, A., Chan, B. Y., Markovich, S., and Wan, X. (2023). Don’t let MEV slip: The costs of swapping on the Uniswap protocol. *arXiv:2309.13648*.
- [20] Wu, W. and Liu, C. (2026). Stability anchors and risk amplifiers: Tail spillovers across stablecoin designs. *arXiv:2602.18820*.
- [21] Hernandez Cruz, W., Xu, J., Tasca, P., and Campajola, C. (2024). No questions asked: Effects of transparency on stablecoin liquidity during the collapse of Silicon Valley Bank. *arXiv:2407.11716*.
- [22] Mohl, V., Frey, S., Leyland, R., Li, K., Nigmatulin, G., Cucuringu, M., Zohren, S., Foerster, J., and Calinescu, A. (2025). JaxMARL-HFT: GPU-accelerated large-scale multi-agent reinforcement learning for high-frequency trading. In *ICAIF '25*. *arXiv:2511.02136*.
- [23] Olby, O., Bacalum, A., Baggott, R., and Stillman, N. (2025). Right place, right time: Market simulation-based RL for execution optimisation. In *ICAIF '25*. *arXiv:2510.22206*.